



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich



# Introduction to the Summer School MiniApp

Sebastian Keller, Javier Otero, Prashanth Kanduri  
and Ben Cumming, CSCS

July 21, 2021

# Overview

In this session we will cover:

1. What is a miniapp?
2. The summer school miniapp overview.
3. First look at the code.
4. Compile, run and visualize the miniapp.

# What is a HPC miniapp?

- Full HPC applications are complicated.
  - Difficult to model/understand performance behavior.
- A miniapp is a smaller code that aim to characterize performance of larger applications.
  - simpler to understand and benchmark than full applications.
  - can be used to test different hardware, languages and libraries.
  - good for learning new techniques!

# The Summer School Miniapp

- Throughout the summer school we will be using a miniapp to reinforce the lessons.
  - During talks there will be small programming exercises to test out what you learn.
  - Then you will get the opportunity to apply the techniques to the miniapp.
- We will start with a serial version that has no parallel optimizations.
- By the end of the course we will have different versions, one for each technique.


# The Application

- The code solves **Fisher's equation**, a **reaction diffusion** model:

$$\frac{\partial s}{\partial t} = D \left( \frac{\partial^2 s}{\partial x^2} + \frac{\partial^2 s}{\partial y^2} \right) + Rs(1 - s).$$

- Used to simulate travelling waves and simple population dynamics.
  - The species  $s$  diffuses.
  - The species reproduces to a maximum of  $s = 1$ .

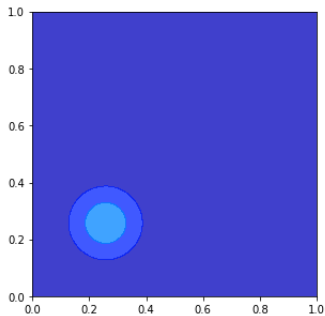
# Initial and Boundary Conditions



The domain is rectangular, with fixed value of  $s = 0$  on each boundary, and a circular region of  $s = 0.1$  in the lower left corner initially.


# Time Evolution of the Solution

$$t = 0.001$$



# Time Evolution of the Solution


$$t = 0.005$$






# Time Evolution of the Solution

$$t = 0.01$$



# Time Evolution of the Solution

$$t = 0.02$$



# Numerical Solution

- The rectangular domain is discretized with a grid of dimension  $nx \times ny$  points.
- A finite volume discretization and method of lines gives the following ordinary differential equation for each grid point

$$\frac{ds_{i,j}}{dt} = \frac{D}{\Delta x^2} (-4s_{i,j} + s_{i+1,j} + s_{i-1,j} + s_{i,j+1} + s_{i,j-1}) + Rs_{i,j}(1 - s_{i,j})$$

$$f_{ij} = [-(4 + \alpha)s_{ij} + s_{i+1,j} + s_{i-1,j} + s_{i,j+1} + s_{i,j-1}]^{k+1} + \alpha s_{ij}^k = 0$$

# Numeric Solution

- One nonlinear equation for each grid point:
  - together they form a system of  $N = nx \times ny$  equations
  - solve with Newton's method
- Each iteration of Newton's method solves a linear system
  - use a matrix-free Conjugate Gradient solver
- Solve the nonlinear system at each time step
  - requires in the order of between 5–10 conjugate gradient iterations

- Don't worry if you don't understand everything.
- We don't need a deep understanding of the mathematics or domain problem to optimize the code.
  - I often work on codes with little domain knowledge.
- The miniapp has a handful of kernels that can be parallelized.
- And care was taken when designing it to make parallelization as easy as possible.
- So let's look a little closer at each part of the code...

# The Code

- The application is written in C++.
- It could be faster...
  - We avoid aggressive optimization to make the code easier to understand.
  - It is not a fine example of design.

# Code Walkthrough

There are three main files of interest:

1. `main.cpp`: Initialization and time stepping code.
2. `linalg.cpp`: BLAS level-1 vector-vector operations and conjugate gradient solver.
3. `operators.cpp` The stencil kernel.

The vector-vector kernels and diffusion operator are the only kernels that have to be parallelized.

# Linear Algebra: linalg.cpp

- This file defines simple kernels for operating on vectors, e.g.:
  - dot product  $x^T y$  or  $x \cdot y$ : `ss_dot`.
  - linear combination  $z = \alpha x + \beta y$ : `ss_lcomb`.
- The kernels of interest are named `ss_XXXX`.
- Each will have to be parallelized using CUDA, MPI and OpenACC.
- The `ss_cg` function implements conjugate gradient using the vector and stencil operations.




# Stencil operator: operators.cpp

This file has the function that applies the stencil operator:

```
for i=2:nx-1
    for j=2:ny-1
        S(i,j) = -(4. + alpha) * U(i,j)
                    + U(i-1,j) + U(i+1,j)
                    + U(i,j-1) + U(i,j+1)
                    + alpha * x_old(i,j)
                    + dxs * U(i,j) * (1.0 - U(i,j));
    end
end
```


# Stencil operator: Interior grid points



interior points have all  
neighbours available

$$S(i,j) = -(4+\alpha)*U(i,j) + U(i-1,j) + U(i+1,j) + U(i,j-1) + U(i,j+1) + \dots$$

# Stencil operator: Boundary grid points



Points on the boundary need to use one or two external boundary points.

$$S(i,j) = -(4+\alpha)*U(i,j) + U(i-1,j) + \text{bndE}[j] + U(i,j-1) + U(i,j+1) + \dots$$

# Testing the Code

Get the code and compile miniapp

```
> git clone git<at>github.com:eth-cscs/SummerSchool2020.git
> cd SummerSchool2020/miniapp/openmp
> module load daint-gpu
> module swap PrgEnv-cray PrgEnv-gnu
> make
```

Run the miniapp

```
> srun -Cgpu --reservation=course ./main 128 128 100 0.01
=====
Welcome to mini-stencil!
version      :: C++  serial
mesh         :: 128 * 128 dx = 0.00787402
time         :: 128 time steps from 0 .. 0.01
iteration    :: CG 200, Newton 50, tolerance 1e-06
=====
-----
simulation took 1.07502 seconds
7439 conjugate gradient iterations, at rate of 6919.88 iters
      /second
959 newton iterations
-----
```

# Exercise: run the miniapp

- Run with 4 different resolutions

- |     |     |     |      |
|-----|-----|-----|------|
| 128 | 128 | 100 | 0.01 |
|-----|-----|-----|------|
- |     |     |     |      |
|-----|-----|-----|------|
| 256 | 256 | 200 | 0.01 |
|-----|-----|-----|------|
- |     |     |     |      |
|-----|-----|-----|------|
| 512 | 512 | 200 | 0.01 |
|-----|-----|-----|------|
- |      |      |     |      |
|------|------|-----|------|
| 1024 | 1024 | 400 | 0.01 |
|------|------|-----|------|

- For each case record:
  1. the number of CG iterations.
  2. the number of CG iterations per second.
- We will refer to these results when testing the MPI and GPU versions of the code.

## Exercise: visualize the results

- The application generates two data files with the final solution: `output.bin` and `output.bov`.
- There is a Python script that will show a contour plot of the solution.
- Now is a good time to test if X-windows is working.

```
> module load daint-gpu
> module load PyExtensions/python3-CrayGNU-20.11
> python3 ./plotting.py -s # -s to get image in pop up
```



**CSCS**

Centro Svizzero di Calcolo Scientifico  
Swiss National Supercomputing Centre

**ETH** zürich

# Questions?

---